

Lecture 6: Sorting Lower Bound and “Linear-Time” Sorting

Jessica Sorrell

September 11, 2025

601.433/633 Introduction to Algorithms

Slides by Michael Dinitz

Introduction

Lots of ways of sorting in $O(n \log n)$ time: mergesort, heapsort, randomized quicksort, deterministic quicksort with BPFRT pivot selection, ...

Is it possible to do better?

Introduction

Lots of ways of sorting in $O(n \log n)$ time: mergesort, heapsort, randomized quicksort, deterministic quicksort with BPFRT pivot selection, ...

Is it possible to do better? No!

Introduction

Lots of ways of sorting in $O(n \log n)$ time: mergesort, heapsort, randomized quicksort, deterministic quicksort with BPFRT pivot selection, ...

Is it possible to do better? No! And yes!

Introduction

Lots of ways of sorting in $O(n \log n)$ time: mergesort, heapsort, randomized quicksort, deterministic quicksort with BPFRT pivot selection, ...

Is it possible to do better? No! And yes!

Comparison Model: we are given a constant-time algorithm which can compare any two elements. No other information about elements.

- ▶ All algorithms we've seen so far have been in this model

Introduction

Lots of ways of sorting in $O(n \log n)$ time: mergesort, heapsort, randomized quicksort, deterministic quicksort with BPFRT pivot selection, ...

Is it possible to do better? No! And yes!

Comparison Model: we are given a constant-time algorithm which can compare any two elements. No other information about elements.

- ▶ All algorithms we've seen so far have been in this model

No: every algorithm in the comparison model must have worst-case running time $\Omega(n \log n)$.

Yes: If we assume extra structure for the elements, can do sorting in $O(n)$ time*

Sorting Lower Bound

Statement

Theorem

Any sorting algorithm in the comparison model must make at least $\log(n!) = \Theta(n \log n)$ comparisons (in the worst case).

Lower bound on the number of comparisons – running time could be even worse!

Allows algorithm to reorder elements, copy them, move them, etc. for free.

$$n! = n \cdot n-1 \cdot \dots \cdot 2 \cdot 1$$
$$\log(n!) = \sum_{i=0}^{n-1} \log(n-i)$$

Statement

Theorem

Any sorting algorithm in the comparison model must make at least $\log(n!) = \Theta(n \log n)$ comparisons (in the worst case).

Lower bound on the number of comparisons – running time could be even worse!

Allows algorithm to reorder elements, copy them, move them, etc. for free.

Why is this hard?

- ▶ Lower bound needs to hold for *all* algorithms
- ▶ How can we simultaneously reason about algorithms as different as mergesort, quicksort, heapsort, ...?

Sorting as Permutations

Think of an array \mathbf{A} as a *permutation*: $\mathbf{A}[i]$ is the $\pi(i)$ 'th smallest element

$$\mathbf{A} = [23, 14, 2, 5, 76]$$

Handwritten: $\mathbf{A}' = [\underline{2}, \underline{5}, \underline{14}, \underline{23}, \underline{76}]$

Corresponds to $\pi = (\underline{3}, \underline{2}, \underline{0}, \underline{1}, \underline{4})$:

$$\underline{\pi(0) = 3}$$

$$\underline{\pi(1) = 2}$$

$$\underline{\pi(2) = 0}$$

$$\underline{\pi(3) = 1}$$

$$\underline{\pi(4) = 4}$$

Sorting as Permutations

Think of an array \mathbf{A} as a *permutation*: $\mathbf{A}[i]$ is the $\pi(i)$ 'th smallest element

$$\mathbf{A} = [23, 14, 2, 5, 76]$$

Corresponds to $\pi = (3, 2, 0, 1, 4)$:

$$\pi(0) = 3$$

$$\pi(1) = 2$$

$$\pi(2) = 0$$

$$\pi(3) = 1$$

$$\pi(4) = 4$$

Lemma

Given \mathbf{A} with $|\mathbf{A}| = n$, if can sort in $T(n)$ comparisons then can find π in $T(n)$ comparisons

Sorting As Permutations (cont'd)

Lemma

Given \mathbf{A} with $|\mathbf{A}| = n$, if can sort in $T(n)$ comparisons then can find π in $T(n)$ comparisons

Proof Sketch.

- ▶ “Tag” each element of \mathbf{A} with index:
 $[23, 14, 2, 5, 76] \rightarrow [(23, 0), (14, 1), (2, 2), (5, 3), (76, 4)]$
- ▶ Sort tagged \mathbf{A} into tagged \mathbf{B} with $T(n)$ comparisons:
 $[(2, 2), (5, 3), (14, 1), (23, 0), (76, 4)]$
- ▶ Iterate through to get π : $\pi(2) = 0, \pi(3) = 1, \pi(1) = 2, \pi(0) = 3, \pi(4) = 4$ □

Sorting As Permutations (cont'd)

Lemma

Given \mathbf{A} with $|\mathbf{A}| = n$, if can sort in $T(n)$ comparisons then can find π in $T(n)$ comparisons

Proof Sketch.

- ▶ “Tag” each element of \mathbf{A} with index:
 $[23, 14, 2, 5, 76] \rightarrow [(23, 0), (14, 1), (2, 2), (5, 3), (76, 4)]$
- ▶ Sort tagged \mathbf{A} into tagged \mathbf{B} with $T(n)$ comparisons:
 $[(2, 2), (5, 3), (14, 1), (23, 0), (76, 4)]$
- ▶ Iterate through to get π : $\pi(2) = 0, \pi(3) = 1, \pi(1) = 2, \pi(0) = 3, \pi(4) = 4$ □

Corollary

If need at least $T(n)$ comparisons to find π , need at least $T(n)$ comparisons to sort!

Generic Algorithm

Want to show that it takes $\Omega(n \log n)$ comparisons to find π in comparison model.

- ▶ Only comparisons cost us anything!

Generic Algorithm

Want to show that it takes $\Omega(n \log n)$ comparisons to find π in comparison model.

- ▶ Only comparisons cost us anything!

Arbitrary algorithm:

- ▶ Starts with some comparison (e.g., compares $A[0]$ to $A[1]$)
- ▶ Rules out some possible permutations!
 - ▶ If $A[0] < A[1]$ then $\pi(0) < \pi(1)$
 - ▶ If $A[0] > A[1]$ then $\pi(1) > \pi(0)$
- ▶ Depending on outcome, choose next comparison to make.
- ▶ Continue until only one possible permutation.

Generic Algorithm

Want to show that it takes $\Omega(n \log n)$ comparisons to find π in comparison model.

- ▶ Only comparisons cost us anything!

Arbitrary algorithm:

- ▶ Starts with some comparison (e.g., compares $A[0]$ to $A[1]$)
- ▶ Rules out some possible permutations!
 - ▶ If $A[0] < A[1]$ then $\pi(0) < \pi(1)$
 - ▶ If $A[0] > A[1]$ then $\pi(1) > \pi(0)$
- ▶ Depending on outcome, choose next comparison to make.
- ▶ Continue until only one possible permutation.

Remind you of anything?

Decision Trees

Model any algorithm as a *binary decision tree*


- ▶ Internal nodes: comparisons
- ▶ Leaves: permutations

Decision Trees

Model any algorithm as a *binary decision tree*

- ▶ Internal nodes: comparisons
- ▶ Leaves: permutations

Example: $n = 3$. Six possible permutations.




Decision Trees

Model any algorithm as a *binary decision tree*

- ▶ Internal nodes: comparisons
- ▶ Leaves: permutations

Example: $n = 3$. Six possible permutations.




Max # comparisons:

Decision Trees

Model any algorithm as a *binary decision tree*


- ▶ Internal nodes: comparisons
- ▶ Leaves: permutations

Example: $n = 3$. Six possible permutations.




Max # comparisons: **3**

Finishing Up



Scale to general n . Consider arbitrary decision tree.


Finishing Up



Scale to general n . Consider arbitrary decision tree.

Max # comparisons


Finishing Up



Scale to general n . Consider arbitrary decision tree.

Max # comparisons = depth of tree


Finishing Up



Scale to general n . Consider arbitrary decision tree.

Max # comparisons = depth of tree
 $\geq \log_2(\# \text{ leaves})$


Finishing Up



Scale to general n . Consider arbitrary decision tree.

Max # comparisons = depth of tree
 $\geq \log_2(\# \text{ leaves})$
 $= \log_2(n!)$

Finishing Up



Scale to general n . Consider arbitrary decision tree.

$$\begin{aligned}\text{Max \# comparisons} &= \text{depth of tree} \\ &\geq \log_2(\# \text{ leaves}) \\ &= \log_2(n!) \\ &= \Theta(n \log n)\end{aligned}$$

Sorting Lower Bound Summary

Theorem

Every sorting algorithm in the comparison model must make at least $\log(n!) = \Theta(n \log n)$ comparisons (in the worst case).

Proof Sketch.

1. Lower bound on finding permutation $\pi \implies$ lower bound on sorting
 2. Any algorithm for finding π is a binary decision tree with $n!$ leaves.
 3. Any binary decision tree with $n!$ leaves has depth $\geq \log(n!) = \Theta(n \log n)$
- \implies Every algorithm has worst case number of comparisons at least $\Theta(n \log n)$. □

“Linear-Time” Sorting

Bypassing the Lower Bound

What if we're *not* in the comparison model?

- ▶ Can do more than just compare elements.

Main example: *integers*.

- ▶ What is the **3**rd bit of **$A[0]$** ?
- ▶ Is **$A[0]$** even?

Same ideas apply to letters, strings, etc.

Counting Sort


Suppose \mathbf{A} consists of n integers, all in $\{0, 1, \dots, k - 1\}$.

Counting Sort


Suppose \mathbf{A} consists of n integers, all in $\{0, 1, \dots, k-1\}$.

Counting Sort:


- ▶ Maintain an array \mathbf{B} of length k initialized to all 0
- ▶ Scan through \mathbf{A} and increment $\mathbf{B}[\mathbf{A}[i]]$.
- ▶ Scan through \mathbf{B} , output i exactly $\mathbf{B}[i]$ times.



$$k-1 = 5$$



12



Counting Sort

Suppose \mathbf{A} consists of n integers, all in $\{0, 1, \dots, k - 1\}$.

Counting Sort:

- ▶ Maintain an array \mathbf{B} of length k initialized to all 0
- ▶ Scan through \mathbf{A} and increment $\mathbf{B}[\mathbf{A}[i]]$.
- ▶ Scan through \mathbf{B} , output i exactly $\mathbf{B}[i]$ times.

Correctness: Obvious

Counting Sort

Suppose \mathbf{A} consists of n integers, all in $\{0, 1, \dots, k - 1\}$.

Counting Sort:

- ▶ Maintain an array \mathbf{B} of length k initialized to all 0
- ▶ Scan through \mathbf{A} and increment $\mathbf{B}[\mathbf{A}[i]]$.
- ▶ Scan through \mathbf{B} , output i exactly $\mathbf{B}[i]$ times.

Correctness: Obvious

Running time:

Counting Sort

Suppose \mathbf{A} consists of n integers, all in $\{0, 1, \dots, k - 1\}$.

Counting Sort:

- ▶ Maintain an array \mathbf{B} of length k initialized to all 0
- ▶ Scan through \mathbf{A} and increment $\mathbf{B}[\mathbf{A}[i]]$.
- ▶ Scan through \mathbf{B} , output i exactly $\mathbf{B}[i]$ times.

Correctness: Obvious

Running time: $O(n + k)$

Bucket Sort: Counting Sort++

Often want to sort *objects* based on *keys*:

- ▶ Each object has a key: integer in $\{0, 1, \dots, k - 1\}$
- ▶ ***A*** consists of ***n*** objects


Bucket Sort: Counting Sort++

Often want to sort *objects* based on *keys*:

- ▶ Each object has a key: integer in $\{0, 1, \dots, k-1\}$
- ▶ **A** consists of **n** objects

Bucket Sort:

- ▶ Same idea as counting sort, but **B**[*i*] is bucket of objects with key *i*
- ▶ Bucket is a linked list with pointers to beginning and end
- ▶ Insert at *end* of list, using end pointer.
- ▶ For output, go through each bucket in order.



Bucket Sort: Counting Sort++

Often want to sort *objects* based on *keys*:

- ▶ Each object has a key: integer in $\{0, 1, \dots, k - 1\}$
- ▶ **A** consists of **n** objects

Bucket Sort:

- ▶ Same idea as counting sort, but **B**[*i*] is bucket of objects with key *i*
- ▶ Bucket is a linked list with pointers to beginning and end
- ▶ Insert at *end* of list, using end pointer.
- ▶ For output, go through each bucket in order.

Running time:

Bucket Sort: Counting Sort++

Often want to sort *objects* based on *keys*:

- ▶ Each object has a key: integer in $\{0, 1, \dots, k - 1\}$
- ▶ **A** consists of **n** objects

Bucket Sort:

- ▶ Same idea as counting sort, but **B**[*i*] is bucket of objects with key *i*
- ▶ Bucket is a linked list with pointers to beginning and end
- ▶ Insert at *end* of list, using end pointer.
- ▶ For output, go through each bucket in order.

Running time: **$O(n + k)$**

Bucket Sort: Counting Sort++

Often want to sort *objects* based on *keys*:

- ▶ Each object has a key: integer in $\{0, 1, \dots, k - 1\}$
- ▶ **A** consists of **n** objects

Bucket Sort:

- ▶ Same idea as counting sort, but **B[i]** is bucket of objects with key **i**
- ▶ Bucket is a linked list with pointers to beginning and end
- ▶ Insert at *end* of list, using end pointer.
- ▶ For output, go through each bucket in order.

Running time: **$O(n + k)$**

Stable: if two objects have same key, order between them after sorting is same as before.

Radix Sort: Setup

What if k is much larger than n , e.g., $k = \Theta(n^2)$?

Radix Sort: Setup

What if k is much larger than n , e.g., $k = \Theta(n^2)$?

Radix sort: $O(n)$ time* for this case

Radix Sort: Setup

What if k is much larger than n , e.g., $k = \Theta(n^2)$?

Radix sort: $O(n)$ time* for this case

Setup:

- ▶ Numbers represented base 10 for historical reasons (all works fine in binary)
- ▶ Assume all numbers have exactly d digits (for simplicity)

Radix Sort: Setup

What if k is much larger than n , e.g., $k = \Theta(n^2)$?

Radix sort: $O(n)$ time* for this case

Setup:

- ▶ Numbers represented base 10 for historical reasons (all works fine in binary)
- ▶ Assume all numbers have exactly d digits (for simplicity)


If you were sorting cards, with a number on each card, what might you do?

Radix Sort: Algorithm

Divide into **10** buckets by first digit, recurse on each bucket by second-digit, etc.


Radix Sort: Algorithm

Divide into **10** buckets by first digit, recurse on each bucket by second-digit, etc.



Radix Sort: Algorithm

Divide into **10** buckets by first digit, recurse on each bucket by second-digit, etc.




Works, but clunky

Radix-Sort: Algorithm (II)

More elegant (and surprising): one bucket, sorting from *least* significant digit to *most*!


Radix-Sort: Algorithm (II)

More elegant (and surprising): one bucket, sorting from *least* significant digit to *most*!



Radix-Sort: Algorithm (II)


More elegant (and surprising): one bucket, sorting from *least* significant digit to *most*!



For iteration i , use bucket sort where key is i 'th digit and object is number.

Radix-Sort: Algorithm (II)

More elegant (and surprising): one bucket, sorting from *least* significant digit to *most*!



For iteration i , use bucket sort where key is i 'th digit and object is number.

Theorem

Radix sort from least significant to most significant is correct if the sort used on each digit is stable.

Least-Significant Radix Sort: Correctness

Proof.

Claim: After i 'th iteration, correctly sorted by last i digits (interpreted as $\#$ in $[0, 10^i - 1]$).

Least-Significant Radix Sort: Correctness

Proof.

Claim: After i 'th iteration, correctly sorted by last i digits (interpreted as $\#$ in $[0, 10^i - 1]$).
Induction on i .

Least-Significant Radix Sort: Correctness

Proof.

Claim: After i 'th iteration, correctly sorted by last i digits (interpreted as $\#$ in $[0, 10^i - 1]$).
Induction on i .

Base case: After first iteration, correctly sorted by last digit

Least-Significant Radix Sort: Correctness

Proof.

Claim: After i 'th iteration, correctly sorted by last i digits (interpreted as $\#$ in $[0, 10^i - 1]$).
Induction on i .

Base case: After first iteration, correctly sorted by last digit

Induction:

- ▶ Suppose correct for i
- ▶ After $i + 1$ sort:

Least-Significant Radix Sort: Correctness

Proof.

Claim: After i 'th iteration, correctly sorted by last i digits (interpreted as $\#$ in $[0, 10^i - 1]$).
Induction on i .

Base case: After first iteration, correctly sorted by last digit

Induction:

- ▶ Suppose correct for i
- ▶ After $i + 1$ sort:
 - ▶ If two numbers have different $i + 1$ digits, now correct.
 - ▶ If two number have same $i + 1$ digit, were correct and still correct by stability.



Least-Significant Radix Sort: Running Time

Recall have n numbers, all numbers have d digits.

Least-Significant Radix Sort: Running Time

Recall have n numbers, all numbers have d digits.

bucket sorts:

Least-Significant Radix Sort: Running Time

Recall have n numbers, all numbers have d digits.

bucket sorts: d

Least-Significant Radix Sort: Running Time

Recall have n numbers, all numbers have d digits.

bucket sorts: d

Time per bucket sort:

Least-Significant Radix Sort: Running Time

Recall have n numbers, all numbers have d digits.

bucket sorts: d

Time per bucket sort: $O(n + k) = O(n + 10) = O(n)$.

Least-Significant Radix Sort: Running Time

Recall have n numbers, all numbers have d digits.

bucket sorts: d

Time per bucket sort: $O(n + k) = O(n + 10) = O(n)$.

Total time: $O(dn)$

Least-Significant Radix Sort: Running Time

Recall have n numbers, all numbers have d digits.

bucket sorts: d

Time per bucket sort: $O(n + k) = O(n + 10) = O(n)$.

Total time: $O(dn)$

Is this good? Bad? In between?

If all numbers distinct, $d \geq \log_{10} n \implies$ total time $O(n \log n)$

Least-Significant Radix Sort: Running Time

Recall have n numbers, all numbers have d digits.

bucket sorts: d

Time per bucket sort: $O(n + k) = O(n + 10) = O(n)$.

Total time: $O(dn)$

Is this good? Bad? In between?

If all numbers distinct, $d \geq \log_{10} n \implies$ total time $O(n \log n)$

Bad: not $O(n)$ time!

Good: “Size of input” is $N = nd$, so linear in size of input!

Least-Significant Radix Sort: Running Time

Recall have n numbers, all numbers have d digits.

bucket sorts: d

Time per bucket sort: $O(n + k) = O(n + 10) = O(n)$.

Total time: $O(dn)$

Is this good? Bad? In between?

If all numbers distinct, $d \geq \log_{10} n \implies$ total time $O(n \log n)$

Bad: not $O(n)$ time!

Good: “Size of input” is $N = nd$, so linear in size of input!

Improve to $O(n)$?

Fast Radix Sort

Change to go **b** digits at a time instead of just **1**.

- ▶ Kind of cheating: look at **b** digits in constant time.
- ▶ Necessary if we want time better than **nd**

Fast Radix Sort

Change to go **b** digits at a time instead of just **1**.

- ▶ Kind of cheating: look at **b** digits in constant time.
- ▶ Necessary if we want time better than **nd**

bucket sorts:

Fast Radix Sort

Change to go **b** digits at a time instead of just **1**.

- ▶ Kind of cheating: look at **b** digits in constant time.
- ▶ Necessary if we want time better than **nd**

bucket sorts: **d/b**

Fast Radix Sort

Change to go **b** digits at a time instead of just **1**.

- ▶ Kind of cheating: look at **b** digits in constant time.
- ▶ Necessary if we want time better than **nd**

bucket sorts: **d/b**

Time per bucket sort:

Fast Radix Sort

Change to go **b** digits at a time instead of just **1**.

- ▶ Kind of cheating: look at **b** digits in constant time.
- ▶ Necessary if we want time better than **nd**

bucket sorts: **d/b**

Time per bucket sort: **$O(n + k) = O(n + 10^b)$**

Fast Radix Sort

Change to go **b** digits at a time instead of just **1**.

- ▶ Kind of cheating: look at **b** digits in constant time.
- ▶ Necessary if we want time better than **nd**

bucket sorts: **d/b**

Time per bucket sort: **$O(n + k) = O(n + 10^b)$**

Total time: **$O\left(\frac{d}{b} (n + 10^b)\right)$**

Fast Radix Sort

Change to go **b** digits at a time instead of just **1**.

- ▶ Kind of cheating: look at **b** digits in constant time.
- ▶ Necessary if we want time better than **nd**

bucket sorts: **d/b**

Time per bucket sort: **$O(n + k) = O(n + 10^b)$**

Total time: **$O\left(\frac{d}{b} (n + 10^b)\right)$**

Set **$b = \log_{10} n$** . If **$d = O(\log n)$** , then time

$$O\left(\frac{d}{\log_{10} n} (n + n)\right) = O(n)$$

Fast Radix Sort

Change to go **b** digits at a time instead of just **1**.

- ▶ Kind of cheating: look at **b** digits in constant time.
- ▶ Necessary if we want time better than **nd**

bucket sorts: **d/b**

Time per bucket sort: **$O(n + k) = O(n + 10^b)$**

Total time: **$O\left(\frac{d}{b} (n + 10^b)\right)$**

Set **$b = \log_{10} n$** . If **$d = O(\log n)$** , then time

$$O\left(\frac{d}{\log_{10} n} (n + n)\right) = O(n)$$

Example: sorting integers between **0** and **n^{10}** . Then **d** should be about **$\log_{10} n^{10} = 10 \log_{10} n$** , as required.